


## Announcements

- Due dates:
  - Group Project #2: Fri, Feb 21<sup>st</sup>
  - Individual HW #4: Mon, Feb 24<sup>th</sup>
- Additional XML slides
  - [Lecture notes](#): [[xml-sql.pdf](#)]
  - ... also linked from HW#4
- Today: Wrapping up OLAP
- Monday: President's Day
- Wednesday: Midterm

ECS-165B

1



## On-Line Analytical Processing (OLAP)

### Demonstration


*Please watch at [db-class.org](http://db-class.org)!*

- *How many have done so already?*
- *What about exercises, quizzes, exams?*

## More on Data Warehouses, and OLAP

Based on material from: Torsten Grust, Alfons Kemper

### Star-Schema: Sales Data



# Star-Schema

Sales					
Date	Store	Product	Count	Customer	SalesRep
25-Jul-00	Passau	1347	1	4711	825
...	...	...	...	...	...

• FACT tables (typically VERY LARGE, 1.000.000+ tuples)

Stores			
City	Country	State	...
Passau	D	Bayern	...
...	...	...	...

Customer			
CustId	Name	Age	...
4711	Kemper	43	...
...	...	...	...

Dimension tables (fairly small)

SalesRep					
PersonId	Name	Expertise	Manager	Age	...
825	Handyman	Electronics	119	23	...
...	...	...	...	...	...

Star-Schema (cont'd)							
Time							
Date	Day	Month	Year	Quarter	Week	Weekday	Season
25-Jul-00	25	7	2000	3	30	Tuesday	Summer
...	...	...	...	...	...		
18-Dec-01	18	12	2001	4	52	Tuesday	Christmas
...	...	...	...	...	...	...	...
• Typical size: 1,000 (3 years)							
Products							
ProduktNr	Produkttyp	Produktgruppe	Produkthauptgruppe	Hersteller	..		
1347	Handy	Mobiltelekom	Telekom	Siemens	..		
...	...	...	...	...	..		
• Typical size: 10,000 (catalog)							

## Non-normalized dimension tables: Hierarchical classification

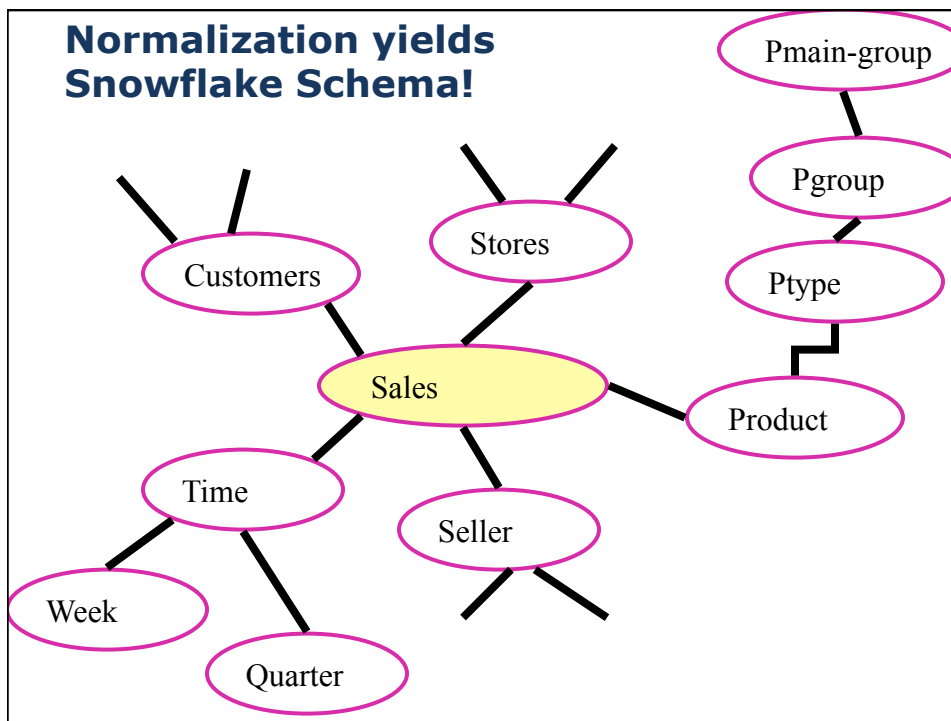
Time							
Date	Day	Month	Year	Quarter	Week	Weekday	Season
25-Jul-00	25	7	2000	3	30	Dienstag	Hochsommer
...	...	...	...	...	...	...	...
18-Dec-01	18	12	2001	4	52	Dienstag	Weihnachten
...	...	...	...	...	...	...	...

FDs: Date → Month → Quarter

Produkte					
ProduktNr	Produkttyp	Produktgruppe	Produkthauptgruppe	Hersteller	..
1347	Handy	Mobiltelekom	Telekom	Siemens	..
...	...	...	...	...	..

ProduktNr → Produkttyp → Produktgruppe → Produkthauptgruppe

## Normalization yields Snowflake Schema!




## Typical Query over a Star Schema: Analysis

How many cell phones of which brand have young customers Ordered during Christmas 2011?



## Algebra-Expression (Star Join)



## Roll-up/Drill-Down

**SELECT** Jahr, Hersteller, **SUM**(v.Anzahl) *quantity*  
**FROM** Verkäufe v, Produkte p, Zeit z  
**WHERE** v.Produkt = p.ProduktNr **AND** v.VerkDatum = z.Datum  
**AND** p.Produkttyp = 'Handy'  
**GROUP BY** p.Hersteller, z.Jahr;  
*Hersteller, Year*

**SELECT** Jahr, **SUM**(v.Anzahl)  
**FROM** Verkäufe v, Produkte p, Zeit z  
**WHERE** v.Produkt = p.ProduktNr **AND** v.VerkDatum = z.Datum  
**AND** p.Produkttyp = 'Handy'  
**GROUP BY** z.Jahr;  
*Year*

**Roll-up** (from detailed to summarized)  
**Drill-down** (from summarized to detailed)

## Analysis along different dimensions

*phone sales*

Hersteller	Jahr	Anzahl
Siemens	1999	2.000
Siemens	2000	3.000
Siemens	2001	3.500
Motorola	1999	1.000
Motorola	2000	1.000
Motorola	2001	1.500
Bosch	1999	500
Bosch	2000	1.000
Bosch	2001	1.500
Nokia	1999	1.000
Nokia	2000	1.500
Nokia	2001	2.000

*dependent attribute measure*  
*Qty*

*{Manufacturer, Year}*

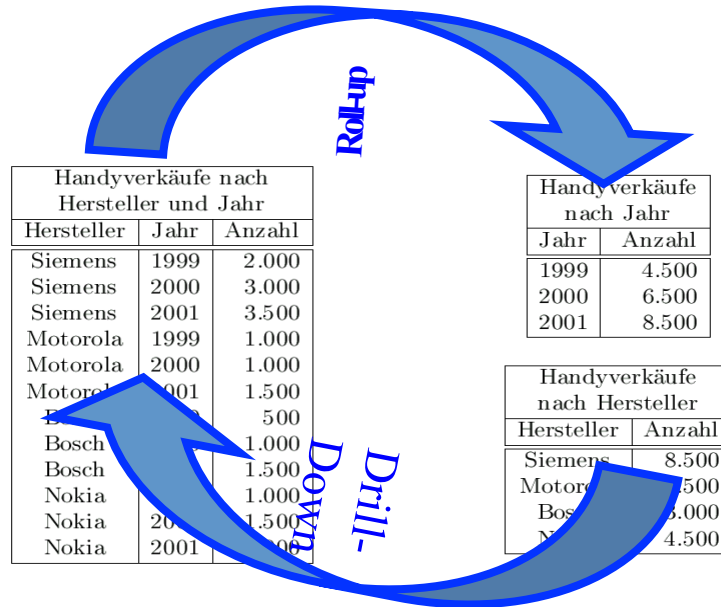
Handyverkäufe nach Jahr	
Jahr	Anzahl
1999	4.500
2000	6.500
2001	8.500

*{Year}*

Handyverkäufe nach Hersteller	
Hersteller	Anzahl
Siemens	8.500
Motorola	3.500
Bosch	3.000
Nokia	4.500

*{Hersteller}*

## Sales Analysis along different dimensions



## Data Cubes ( $n$ -dimensional)

In Decision-Support-Systems similar to spreadsheet-style (*cross tabulation*):

Hersteller \ Jahr	1999	2000	2001	$\Sigma$
Siemens	2.000	3.000	3.500	8.500
Motorola	1.000	1.000	1.500	3.500
Bosch	500	1.000	1.500	3.000
Nokia	1.000	1.500	2.000	4.500
$\Sigma$	4.500	6.500	8.500	19.500

This 2-dimensional **data cube** includes all results from the previous slide

## Extreme case: Maximal Aggregation

No grouping (i.e., no **GROUP BY**-clause)  $\Rightarrow$  all rows form a single group, prior to aggregation:

```
SELECT SUM(Anzahl)
FROM   Verkäufe v, Produkte p
WHERE  v.Produkt = p.ProduktNr AND p.Produkttyp = 'Handy';
```

## Materialization of Aggregates

```
INSERT INTO Handy2DCube
( SELECT p.Hersteller, z.Jahr, SUM(v.Anzahl)
  FROM Verkäufe v, Produkte p, Zeit z
  WHERE v.Produkt = p.ProduktNr and p.Produkttyp = 'Handy'
    and v.VerkDatum = z.Datum
  GROUP BY z.Jahr, p.Hersteller ) UNION
( SELECT p.Hersteller, NULL, SUM(v.Anzahl)
  FROM Verkäufe v, Produkte p
  WHERE v.Produkt = p.ProduktNr and p.Produkttyp = 'Handy'
  GROUP BY p.Hersteller ) UNION
( SELECT NULL, z.Jahr, SUM(v.Anzahl)
  FROM Verkäufe v, Produkte p, Zeit z
  WHERE v.Produkt = p.ProduktNr and p.Produkttyp = 'Handy'
    and v.VerkDatum = z.Datum
  GROUP BY z.Jahr ) UNION
( SELECT NULL, NULL, SUM(v.Anzahl)
  FROM   Verkäufe v, Produkte p
  WHERE v.Produkt = p.ProduktNr and p.Produkttyp = 'Handy' );
```

Attributes = {Year, Maker}

$\Rightarrow$  subqueries for

Group By combination

$\{ \emptyset, \{y, m\}, \{y\}, \{m\} \}$



## Materialization of Aggregates

```

INSERT INTO Handy2DCube
( SELECT p.Hersteller, z.Jahr, SUM(v.Anzahl)
  FROM Verkäufe v, Produkte p, Zeit z
  WHERE v.Produkt = p.ProduktNr and p.Produkttyp = 'Handy'
        and v.VerkDatum = z.Datum
  GROUP BY z.Jahr, p.Hersteller ) UNION {Y,M}
( SELECT p.Hersteller, NULL, SUM(v.Anzahl)
  FROM Verkäufe v, Produkte p
  WHERE v.Produkt = p.ProduktNr and p.Produkttyp = 'Handy'
  GROUP BY p.Hersteller ) UNION {H}
( SELECT NULL, z.Jahr, SUM(v.Anzahl)
  FROM Verkäufe v, Produkte p, Zeit z
  WHERE v.Produkt = p.ProduktNr and p.Produkttyp = 'Handy'
        and v.VerkDatum = z.Datum
  GROUP BY z.Jahr ) UNION {Y}
( SELECT NULL, NULL, SUM(v.Anzahl)
  FROM Verkäufe v, Produkte p
  WHERE v.Produkt = p.ProduktNr and p.Produkttyp = 'Handy' );
  
```

*Handwritten notes: {Y,M}, {H}, {Y}, {} No group by*

## Relational structure of the cube

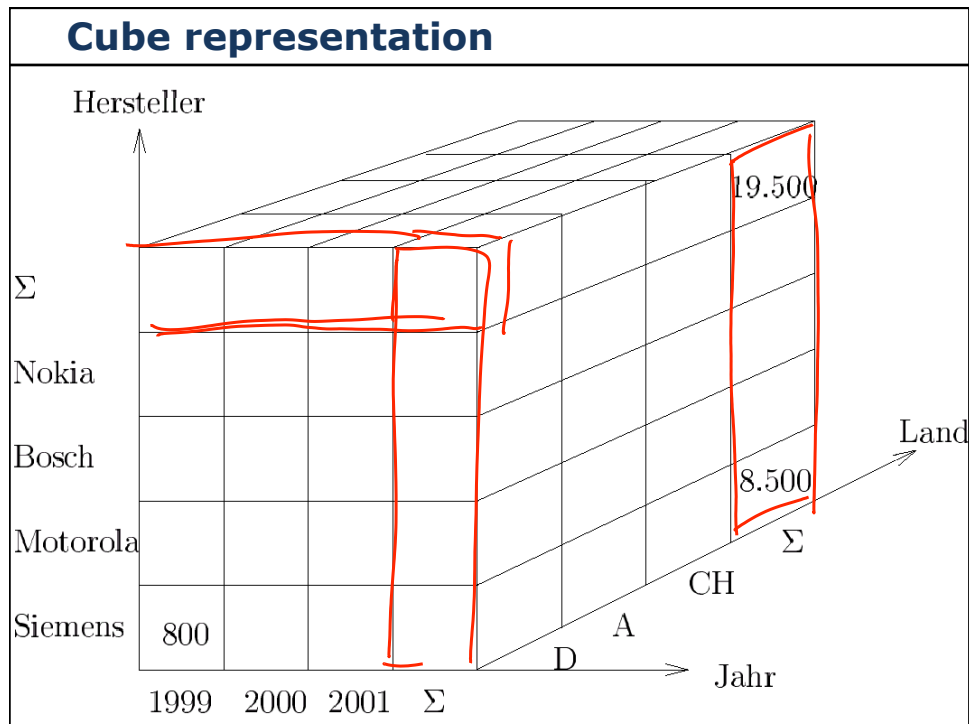
**M Handy2DCube**

Hersteller	Jahr	Anzahl
Siemens	1999	2.000
Siemens	2000	3.000
Siemens	2001	3.500
Motorola	1999	1.000
Motorola	2000	1.000
Motorola	2001	1.500
Bosch	1999	500
Bosch	2000	1.000
Bosch	2001	1.500
Nokia	2000	1.000
Nokia	2001	1.500
Nokia	2001	2.000
<u>null</u>	1999	4.500
<u>null</u>	2000	6.500
<u>null</u>	2001	8.500
Siemens	<u>null</u>	8.500
Motorola	<u>null</u>	3.500
Bosch	<u>null</u>	3.000
Nokia	<u>null</u>	4.500
<u>null</u>	<u>null</u>	19.500

*Handwritten notes: {Y,M}, {H}, {Y}, {} No group by*

**Handy3DCube**

Hersteller	Jahr	Land	Anzahl
Siemens	1999	D	800
Siemens	1999	A	600
Siemens	1999	CH	600
Siemens	2000	D	1.200
Siemens	2000	A	800
Siemens	2000	CH	1.000
Siemens	2001	D	1.400
...	...	...	...
Motorola	1999	D	400
Motorola	1999	A	300
Motorola	1999	CH	300
...	...	...	...
Bosch	...	...	...
...	...	...	...
<u>null</u>	1999	D	...
<u>null</u>	2000	D	...
...	...	...	...
Siemens	<u>null</u>	<u>null</u>	8.500
...	...	...	...
<u>null</u>	<u>null</u>	<u>null</u>	19.500



### The CUBE-Operator (SQL:1999)

- Computing Aggregates leads to hard-to-optimize queries:
  - $n$  Dimensions  $\Rightarrow 2^n$  subqueries, with **UNION**
  - But: aggregate can be computed hierarchically
- This **CUBE**-Operator replaces  $2^3$  subqueries:

```

SELECT p.Hersteller, z.Jahr, f.Land, SUM(v.Anzahl)
FROM   Verkäufe v, Produkte p, Zeit z, Filialen f
WHERE  v.Produkt = p.ProduktNr AND p.Produkttyp = 'Handy'
        AND v.VerkDatum = z.Datum AND v.Filiale = f.Filialenkennung
GROUP BY CUBE (z.Jahr, p.Hersteller, f.Land);
               Year, Maker, Country
  
```

## Reuse of Partial Aggregates

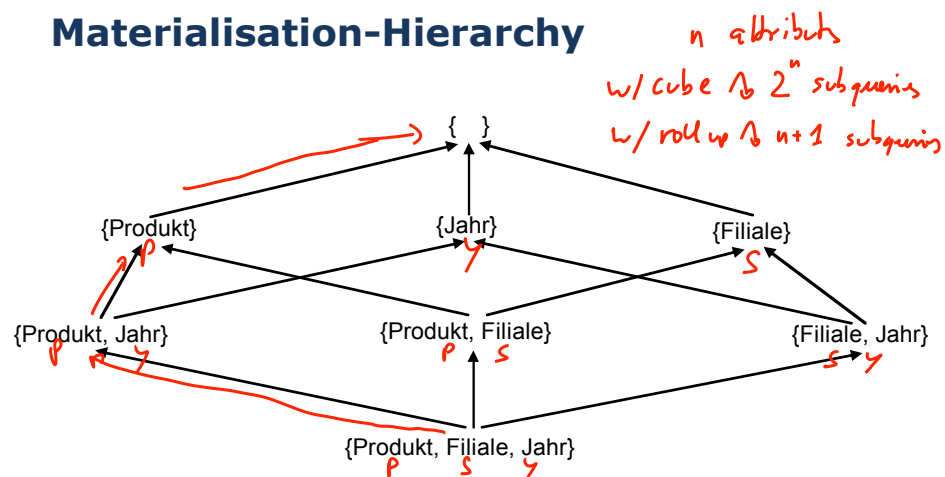
Assume the following aggregate has been materialized:

*Sales Produkt Store Year*  
**INSERT INTO** VerkäufeProduktFilialeJahr  
 ( **SELECT** v.Produkt, v.Filiale, z.Jahr, **SUM**(v.Anzahl)  
**FROM** Verkäufe v, Zeit z  
**WHERE** v.VerkDatum = z.Datum  
**GROUP BY** v.Produkt, v.Filiale, z.Jahr );

Then the following query can use the pre-computed aggregate (instead of the original fact table)

**SELECT** v.Produkt, v.Filiale, **SUM**(v.Anzahl)  
**FROM** Verkäufe v VerkäufeProduktFilialeJahr v  
**GROUP BY** v.Produkt, v.Filiale

## Materialisation-Hierarchy



- Subaggregates S can be **reused** for aggregation A if there is a path  $S \rightarrow \dots \rightarrow A$
- Lattice structure

## WITH ROLLUP vs WITH CUBE

- ROLLUP (Year, Month, Day)

- { Y, M, D }
- { Y, M }
- { Y }
- { }

- CUBE (Year, Month, Day)

- { Y, M, D }
- { Y, M }
- { Y, D }
- { Y }
- { M, D }
- { M }
- { D }
- { }

